

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import datetime as dt
%matplotlib inline
```

1. Data Wrangling

Check each table for values outside of range and missing values.

```
In [50]: df_trans = pd.read_csv('transactions_v2.csv')
```

```
In [6]: df_trans.isnull().values.any()
```

```
Out[6]: False
```

```
In [23]: df_trans.describe()
```

```
Out[23]:
```

	payment_method_id	payment_plan_days	plan_list_price	actual_amount_paid	is_auto_renew	transaction
count	1.431009e+06	1.431009e+06	1.431009e+06	1.431009e+06	1.431009e+06	1.431009e+06
mean	3.791835e+01	6.601770e+01	2.817870e+02	2.813172e+02	7.853025e-01	2.016848e+01
std	4.964805e+00	1.024864e+02	4.351861e+02	4.354200e+02	4.106124e-01	4.858797e+00
min	2.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	2.015010e+00
25%	3.600000e+01	3.000000e+01	9.900000e+01	9.900000e+01	1.000000e+00	2.017023e+01
50%	4.000000e+01	3.000000e+01	1.490000e+02	1.490000e+02	1.000000e+00	2.017031e+01
75%	4.100000e+01	3.000000e+01	1.490000e+02	1.490000e+02	1.000000e+00	2.017032e+01
max	4.100000e+01	4.500000e+02	2.000000e+03	2.000000e+03	1.000000e+00	2.017033e+01

Transactions data table doesn't have any missing values and the ranges of the features look good.

```
In [49]: df_train = pd.read_csv('train_v2.csv')
```

```
In [18]: df_train.isnull().values.any()
```

```
Out[18]: False
```

```
In [20]: df_train.describe()
```

```
Out[20]:
```

	is_churn
count	970960.000000
mean	0.089942
std	0.286099
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
In [4]: df_members = pd.read_csv('members_v3.csv')
```

```
In [39]: df_members.isnull().sum()
```

```
Out[39]: msno                0
city                0
bd                 0
gender            4429505
registered_via      0
registration_init_time  0
dtype: int64
```

The members table has missing gender values.

```
In [23]: df_members['gender'].isnull().sum()/len(df_members)
```

```
Out[23]: 0.65433527838873129
```

There is a large percentage of users whose gender is unknown, 65%. Since the percentage is so large, I will keep the data points with the missing values.

```
In [5]: df_members.describe()
```

```
Out[5]:
```

	city	bd	registered_via	registration_init_time
count	6.769473e+06	6.769473e+06	6.769473e+06	6.769473e+06
mean	3.847358e+00	9.795794e+00	5.253069e+00	2.014518e+07
std	5.478359e+00	1.792590e+01	2.361398e+00	2.318601e+04
min	1.000000e+00	-7.168000e+03	-1.000000e+00	2.004033e+07
25%	1.000000e+00	0.000000e+00	4.000000e+00	2.014042e+07
50%	1.000000e+00	0.000000e+00	4.000000e+00	2.015101e+07
75%	4.000000e+00	2.100000e+01	7.000000e+00	2.016060e+07
max	2.200000e+01	2.016000e+03	1.900000e+01	2.017043e+07

The feature bd gives the users age, and there are some values outside of a range that makes sense. The min age is -7168 and the max age is 2016. The proportion of bad ages is:

```
In [26]: len(df_members[(df_members['bd']<0) | (df_members['bd']>100)])/len(df_members)
```

```
Out[26]: 0.0008347769464476777
```

The max value of registered_via should be 16, and the min value should be 3. However, there are values that lie outside of this range. The proportion of bad values is:

```
In [33]: len(df_members[(df_members['registered_via']<3) | (df_members['registered_via']>16)])/len(df_members)
```

```
Out[33]: 0.000586308564935557
```

Since the number of users with bad ages is a small percentage of the total users, I will remove these rows.

```
In [6]: df_logs = pd.read_csv('user_logs_v2.csv')
```

```
In [28]: df_logs.describe()
```

```
Out[28]:
```

	date	num_25	num_50	num_75	num_985	num_100	num_unq	to
count	1.839636e+07	1.839636e+07	1.839636e+07	1.839636e+07	1.839636e+07	1.839636e+07	1.839636e+07	1.839
mean	2.017032e+07	6.191401e+00	1.508789e+00	9.413759e-01	1.079905e+00	3.028246e+01	2.903615e+01	7.904
std	8.916720e+00	1.342827e+01	3.908539e+00	1.924840e+00	3.518409e+00	4.203641e+01	3.219866e+01	1.013
min	2.017030e+07	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00	1.000
25%	2.017031e+07	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	7.000000e+00	8.000000e+00	1.959
50%	2.017032e+07	2.000000e+00	1.000000e+00	0.000000e+00	0.000000e+00	1.700000e+01	1.800000e+01	4.582
75%	2.017032e+07	7.000000e+00	2.000000e+00	1.000000e+00	1.000000e+00	3.700000e+01	3.800000e+01	9.848
max	2.017033e+07	5.639000e+03	9.120000e+02	5.080000e+02	1.561000e+03	4.110700e+04	4.925000e+03	9.194

```
In [30]: df_logs.isnull().values.any()
```

```
Out[30]: False
```

The logs table doesn't have any null values, and the ranges for the features seems sensible.

2. Data Storytelling

```
In [51]: len(df_train[df_train.is_churn==1])/len(df_train)
```


```
Out[51]: 0.08994191315811156
```

About 10% of the users have churned.

```
In [7]: sns.countplot(x='is_churn',data=df_train)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10288a51d0>
```




The majority of the users did not list their genders. Of the users who did list their gender, the number of male users is almost equal to the number of female users.

```
In [10]: df_members_na = df_members.fillna('NA')
```

```
In [11]: sns.countplot(x='gender',data=df_members_na)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0ffe7185f8>
```




The most popular registration method is method 4. Methods 3 and 9 are the next most popular, followed by method 7, about half as many signups as 3 & 9). The rest of the methods are not very popular.

```
In [13]: sns.countplot(x='registered_via',data=df_members_na)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0ffe4f5be0>
```




Most of the users live in the city labeled 1. The other 21 cities do not have many registrations.

```
In [18]: sns.countplot(x='city',data=df_members_na)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```


```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0ffe0650b8>
```



As we saw before, there are many incorrect values for user age.

```
In [23]: sns.distplot(df_members['bd'])
```


```
Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0ffe0c7550>
```



Plotting only sensible values for the age, shows that most of the users are teenagers and young adults. The distribution of ages peaks around 25, and then decreases.

```
In [26]: sns.distplot(df_members['bd'][(df_members.bd>0) & (df_members.bd<100)])
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0ffdcd3cc0>
```




```
In [5]: df_members['registration_init_time'] = pd.to_datetime(df_members['registration_init_time'],  
format='%Y%m%d')
```

The service wasn't very popular until 2010, at which point signups began to slowly increase. After 2015, signups increased strongly. However, recently new registrations have been declining.

```
In [17]: plt.plot(df_members['registration_init_time'].value_counts().sort_index())
```

```
Out[17]: [matplotlib.lines.Line2D at 0x7f1e69c067f0>]
```




```
In [44]: year = pd.to_datetime(df_members['registration_init_time'], format='%Y%m%d').dt.year  
year = year.to_frame()
```

```
In [46]: sns.countplot(x='registration_init_time',data=year)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1e6763f6d8>
```




```
In [40]: months = pd.to_datetime(df_members['registration_init_time'], format='%Y%m%d').dt.month  
months=months.to_frame()
```

The beginning of the year and the end of the year have higher initial registrations than the other months. The least popular months are April and May. Note that this trend might not hold over all years.

```
In [43]: sns.countplot(x='registration_init_time',data=months)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1e676bb320>
```




```
In [45]: dayofweek = pd.to_datetime(df_members['registration_init_time'], format='%Y%m%d').dt.dayofweek  
dayofweek = dayofweek.to_frame()
```

Weekend days have higher initial registrations than the rest of the week. The weekday with the most registrations is Friday.

```
In [47]: sns.countplot(x='registration_init_time',data=dayofweek)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1e675d7e80>
```



Looking at the mean of is_auto_renew, we see that the majority of users, 78%, have their plans set to automatically renew.

```
In [26]: df_trans.is_auto_renew.mean()
```

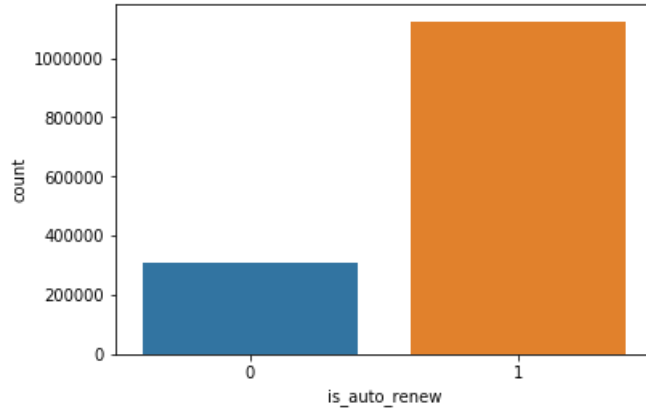
```
Out[26]: 0.7853025382789347
```



```
In [9]: sns.countplot(x='is_auto_renew',data=df_trans)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70e78ed6d8>
```



Only a small proportion of users in the data set canceled their subscriptions. The fraction that cancelled is:


```
In [52]: len(df_trans[df_trans.is_cancel==1])/len(df_trans)
```

```
Out[52]: 0.02455120827332323
```

```
In [11]: sns.countplot(x='is_cancel',data=df_trans)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x7f70e7873128>
```




The most popular payment method is 41.

```
In [53]: sns.countplot(x='payment_method_id',data=df_trans)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```


```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1e6754fe48>
```



```
In [55]: sns.countplot(x='payment_plan_days',data=df_trans)
```

```
/home/rebecca/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py:1460: FutureWarning: remove_na is deprecated and is a private function. Do not use.  
stat_data = remove_na(group_data)
```

```
Out[55]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1df245a358>
```



```
In [ ]:
```

The mean number of unique songs listened to daily is about 29.


```
In [32]: df_logs.num_unq.mean()
```

```
Out[32]: 29.036145516162382
```

All of the auto_renew plans have only 2 lengths (payment_plan_days).

```
In [22]: sns.lmplot(x="payment_method_id", y="payment_plan_days", hue="is_auto_renew", data=df_trans, fit_reg=False)
```


```
Out[22]: <seaborn.axisgrid.FacetGrid at 0x7f70c6f5bd30>
```



These two auto_renew plans have a few different prices.

```
In [28]: sns.lmplot(x="payment_method_id", y="plan_list_price", hue="is_auto_renew", data=df_trans, fit_reg=False)
```


```
Out[28]: <seaborn.axisgrid.FacetGrid at 0x7f70c6ec49e8>
```



From the scatterplot of payment_plan_days vs plan_list_price it looks like there is a max plan price for each plan duration.

```
In [27]: sns.lmplot(x="plan_list_price", y="payment_plan_days", hue="is_auto_renew", data=df_trans, fit_reg=False)
```


```
Out[27]: <seaborn.axisgrid.FacetGrid at 0x7f70c2c166d8>
```



Auto renew is never false when is_cancel is true.

```
In [30]: sns.lmplot(x="is_auto_renew", y="is_cancel", data=df_trans, fit_reg=False)
```

```
Out[30]: <seaborn.axisgrid.FacetGrid at 0x7f70c2baecc0>
```



```
In [19]: df_trans.groupby('msno')['is_auto_renew'].apply(lambda x: x.mode())
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-19-b09305af209d> in <module>()
----> 1 df_trans.groupby('msno')['is_auto_renew'].apply(lambda x: x.mode())

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/groupby.py in apply
(self, func, *args, **kwargs)
    719         # ignore SettingWithCopy here in case the user mutates
    720         with option_context('mode.chained_assignment', None):
--> 721             return self._python_apply_general(f)
    722
    723     def _python_apply_general(self, f):

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/groupby.py in _python_apply_general(self, f)
    723     def _python_apply_general(self, f):
    724         keys, values, mutated = self.grouper.apply(f, self._selected_obj,
--> 725                                                    self.axis)
    726
    727         return self._wrap_applied_output(

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/groupby.py in apply
(self, f, data, axis)
    1878         # group might be modified
    1879         group_axes = _get_axes(group)
-> 1880         res = f(group)
    1881         if not _is_indexed_like(res, group_axes):
    1882             mutated = True

<ipython-input-19-b09305af209d> in <lambda>(x)
----> 1 df_trans.groupby('msno')['is_auto_renew'].apply(lambda x: x.mode())

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/series.py in mode
(self)
    1277         """
    1278         # TODO: Add option for bins like value_counts()
-> 1279         return algorithms.mode(self)
    1280
    1281     @Appender(base._shared_docs['unique'] % _shared_doc_kwargs)

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/algorithms.py in mode
(values)
    675
    676     result = _reconstruct_data(result, original.dtype, original)
--> 677     return Series(result)
    678
    679

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/series.py in __init
__(self, data, index, dtype, name, copy, fastpath)
    262         else:
    263             data = _sanitize_array(data, index, dtype, copy,
--> 264                                   raise_cast_failure=True)
    265
    266             data = SingleBlockManager(data, index, fastpath=True)

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/pandas/core/series.py in _sani
tize_array(data, index, dtype, copy, raise_cast_failure)
    3132
    3133     # GH #846
-> 3134     if isinstance(data, (np.ndarray, Index, Series)):
    3135
    3136         if dtype is not None:
```

```
KeyboardInterrupt:
```

```
In [20]: sns.barplot(x="is_auto_renew", y="is_churn", data=df_trans)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-20-85fcf67b30d4> in <module>()
----> 1 sns.barplot(x="is_auto_renew", y="is_churn", data=df_trans)

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py in barplot(x, y, hue, data, order, hue_order, estimator, ci, n_boot, units, orient, color, palette, saturation, errcolor, errwidth, capsize, dodge, ax, **kwargs)
    2938         estimator, ci, n_boot, units,
    2939         orient, color, palette, saturation,
-> 2940         errcolor, errwidth, capsize, dodge)
    2941
    2942     if ax is None:

~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py in __init__(self, x, y, hue, data, order, hue_order, estimator, ci, n_boot, units, orient, color, palette, saturation, errcolor, errwidth, capsize, dodge)
    1584         """Initialize the plotter."""
    1585         self.establish_variables(x, y, hue, data, orient,
-> 1586                                order, hue_order, units)
    1587         self.establish_colors(color, palette, saturation)
    1588         self.estimate_statistic(estimator, ci, n_boot)

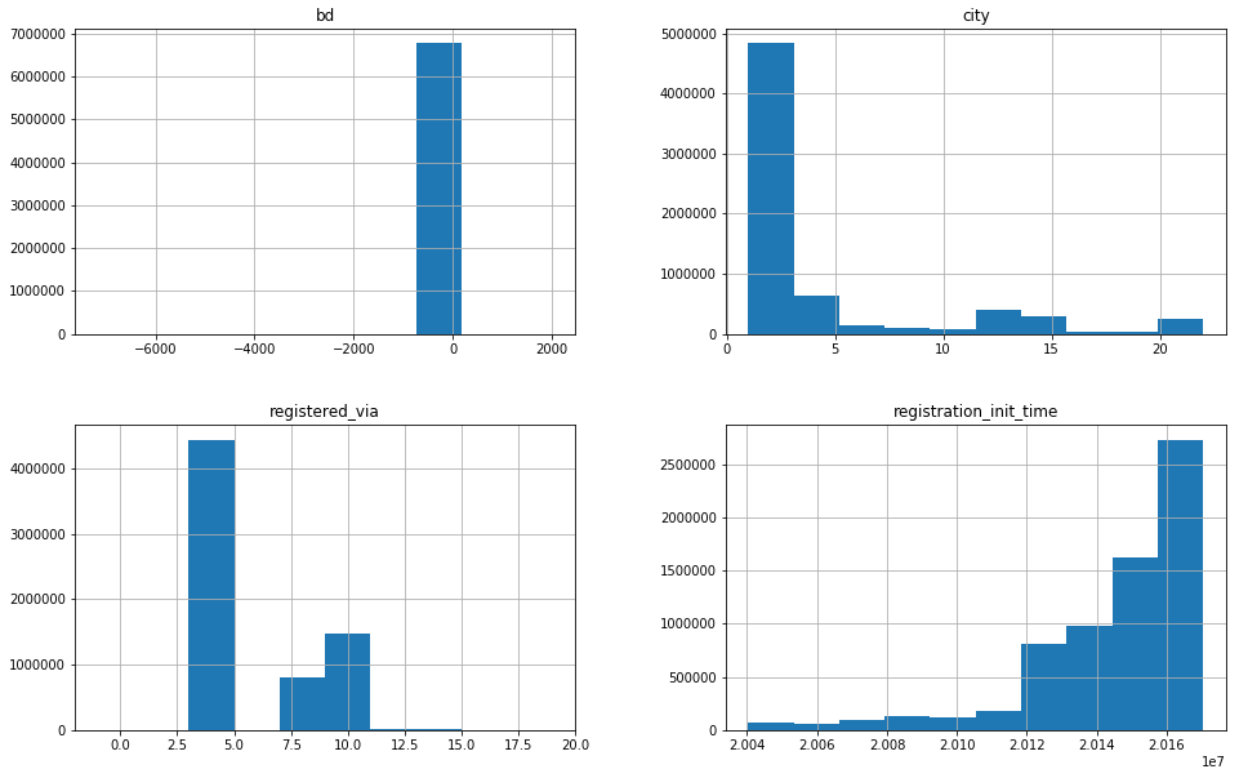
~/anaconda3/envs/my_projects_env/lib/python3.6/site-packages/seaborn/categorical.py in establish_variables(self, x, y, hue, data, orient, order, hue_order, units)
    149         if isinstance(input, string_types):
    150             err = "Could not interpret input '{}'.format(input)
-> 151             raise ValueError(err)
    152
    153         # Figure out the plotting orientation

ValueError: Could not interpret input 'is_churn'
```

```
In [33]: df = df_trans.copy()
df['is_churn'] = df_train['is_churn']
```

```
In [3]: df_members.hist(figsize=(16, 10))
```

```
Out[3]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f0ffedf2898>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f1028a1d710>],  
  [<matplotlib.axes._subplots.AxesSubplot object at 0x7f10289d77f0>,  
  <matplotlib.axes._subplots.AxesSubplot object at 0x7f1028976a20>]], dtype=object)
```



```
In [15]: len(df_members[df_members['bd']>100])
```

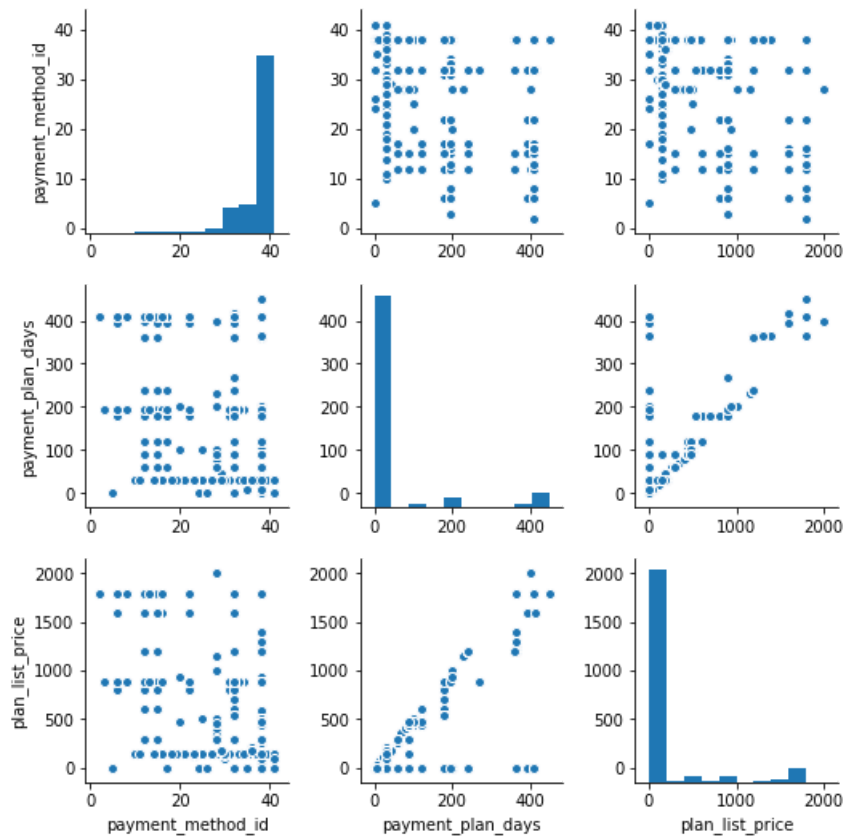
```
Out[15]: 5377
```

gender is missing for some data points. Check to see whether all msno exist in all tables.

```
In [42]: df1 = df_members.copy()  
df1['is_churn'] = df_train['is_churn']
```


```
In [49]: sns.pairplot(df[['payment_method_id', 'payment_plan_days', 'plan_list_price']],diag_kind='hist')
#for ax in g.axes.flat:
#    plt.setp(ax.get_xticklabels(), rotation=45)
```

Out[49]: <seaborn.axisgrid.PairGrid at 0x7f3289dd7cc0>



```
In [51]: sns.pairplot(df[['payment_method_id', 'is_cancel']],diag_kind='hist')
```

Out[51]: <seaborn.axisgrid.PairGrid at 0x7f328134ce48>



```
In [ ]: transactions.hist(column=['actual_amount_paid', 'plan_list_price'],figsize=(16, 5),bins=50)
```

```
In [ ]: gr.plot.line(y=['num_unq'],figsize=(16, 5))
```